

CAPSTONE PROJECT:
GIVE ME SOME
CREDIT

FOUNDATIONS OF DATA SCIENCE WORKSHOP, AUTUMN 2015

DAVID CROOK



GIVE ME SOME CREDIT

KAGGLE.COM

- Credit Scoring modeling competition, originally conducted Autumn 2011. Entries from 925 teams.
- Data contains over 250,000 borrower observations (~100,000 of them for scoring competitor entry), with variables like age, DebtRatio and NumberRealEstateLoansOrLines
- A well-defined *binary classification* problem: Predict a probability on a binary dependent variable.
 - Specifically, whether or not a given borrower will experience "serious financial distress" within two years. 6.6% of 150,000 in training dataset did.

THE APPROACH

ORDER OF OPERATIONS

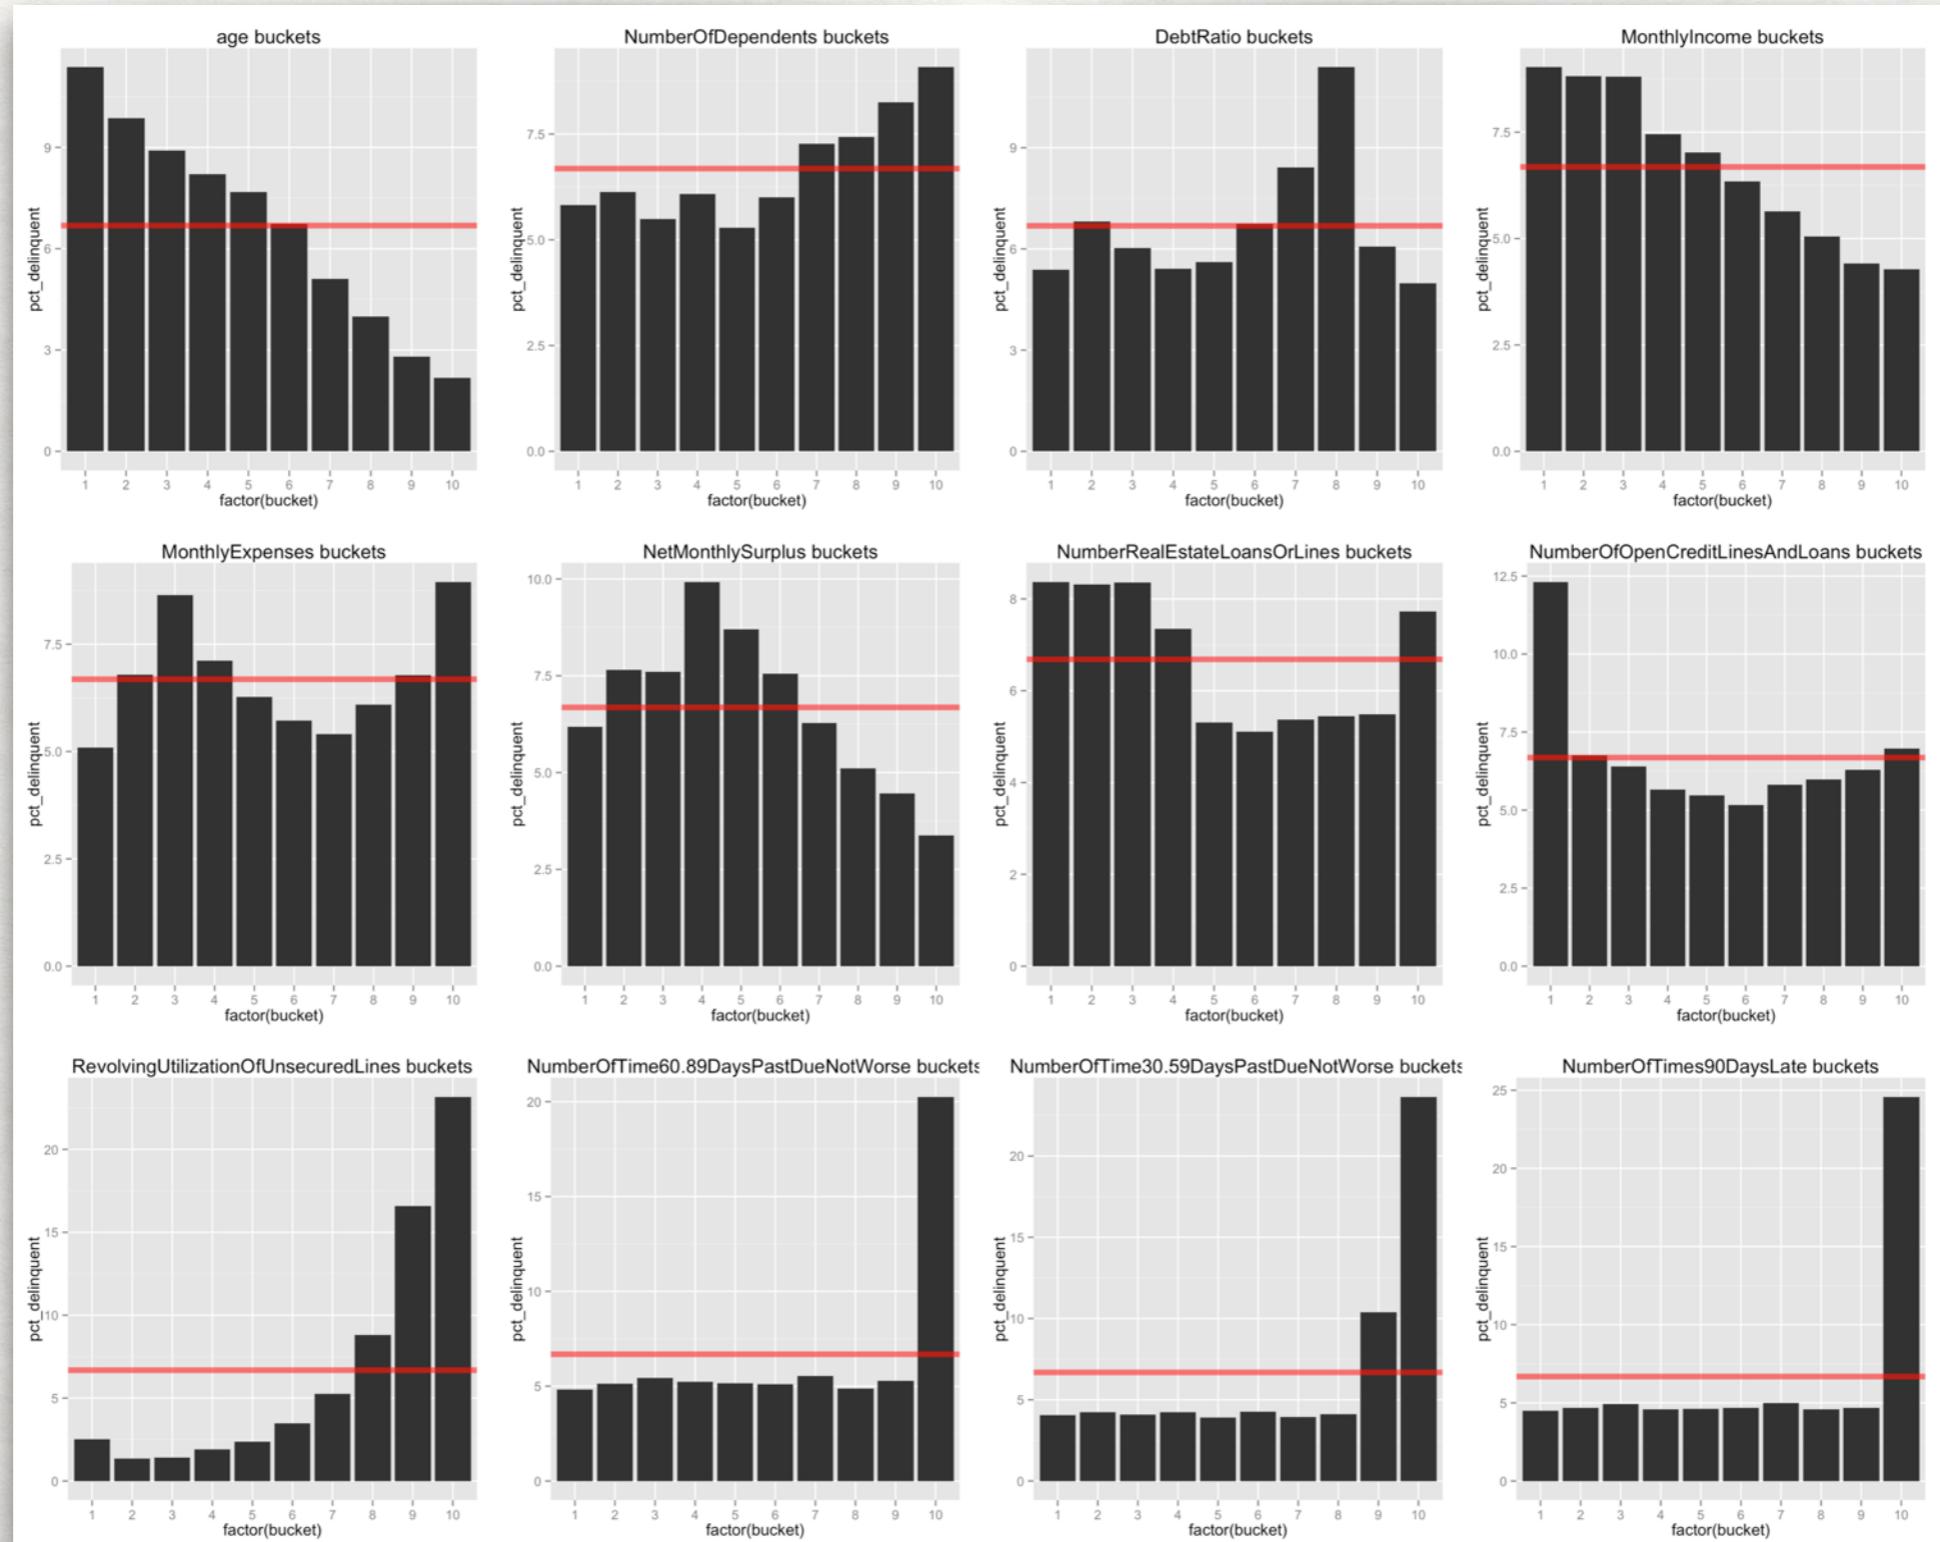
- Define problem (was defined by competition) and Obtain data
- Explore and clean data*
- Build models
- Evaluate + validate models
- Iterate

Refer to paper and code for full details.

* Some Feature engineering was done in this step

TRENDS WITHIN FEATURES

RED LINE: AVERAGE OUTPUT IN ALL OBSERVATIONS



THE APPROACH - BUILD MODELS

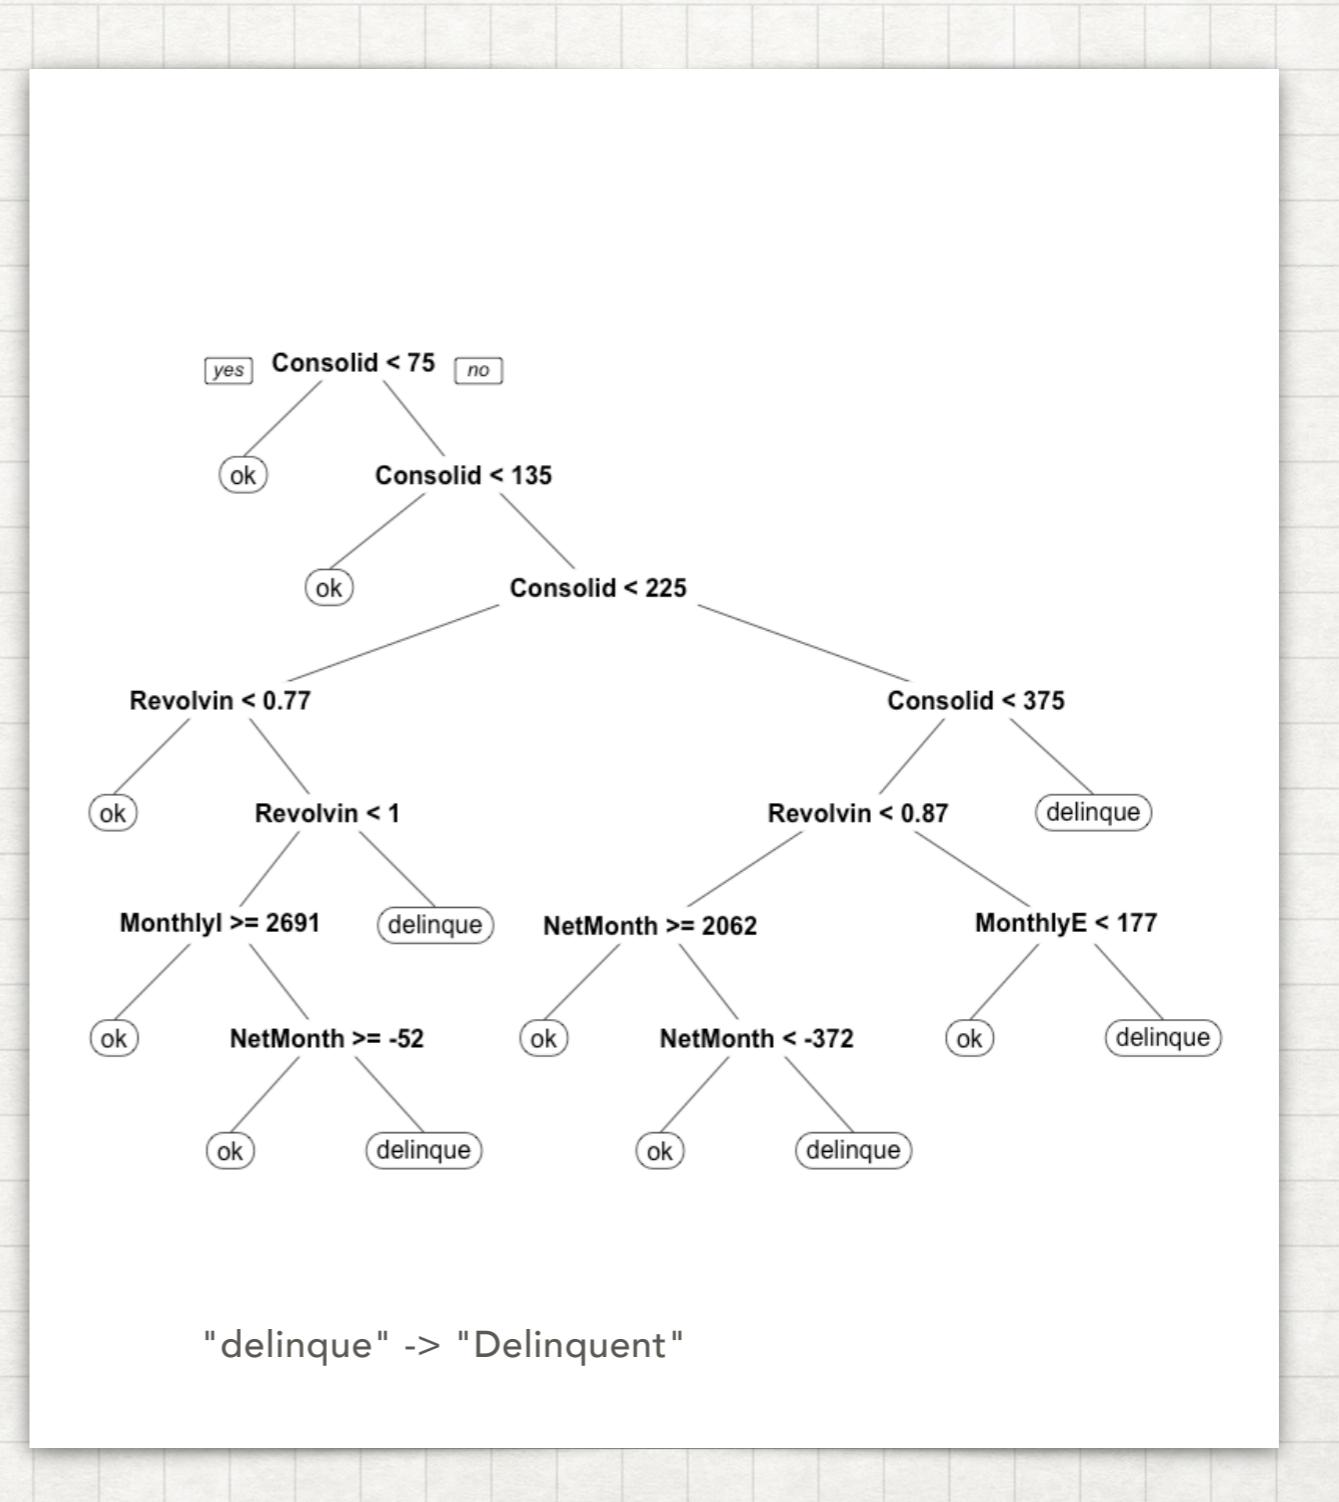
BUILD ML MODELS TRAINED WITH BORROWER DATA

- Three model types built: *Logistic Regression*, *Classification Tree*, and *Random Forest*
- Classification Trees can be more understandable than other types
- Logistic Regressions are binary classifiers
- Random Forest is an ensemble method using classification trees.
 - RF algorithm reduces chance of over-fitting compared to normal classification tree.

EXAMPLE DECISION TREE

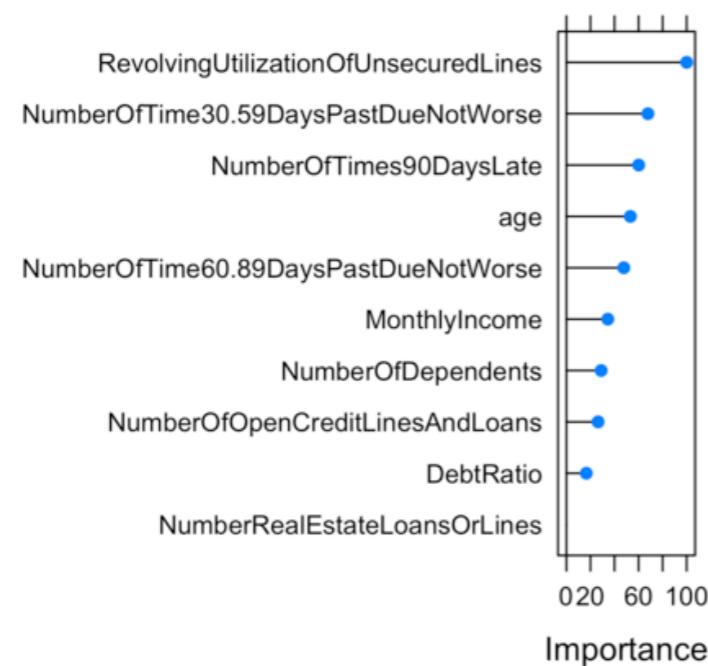
FIVE PREDICTORS USED, BUILT ON FULL TRAINING DATASET

- ConsolidatedNumberOfDaysPastDue - critical feature
 - RevolvingUtilizationOfUnsecuredLines - also very important
 - NetMonthlySurplus
 - MonthlyExpenses
 - MonthlyIncome

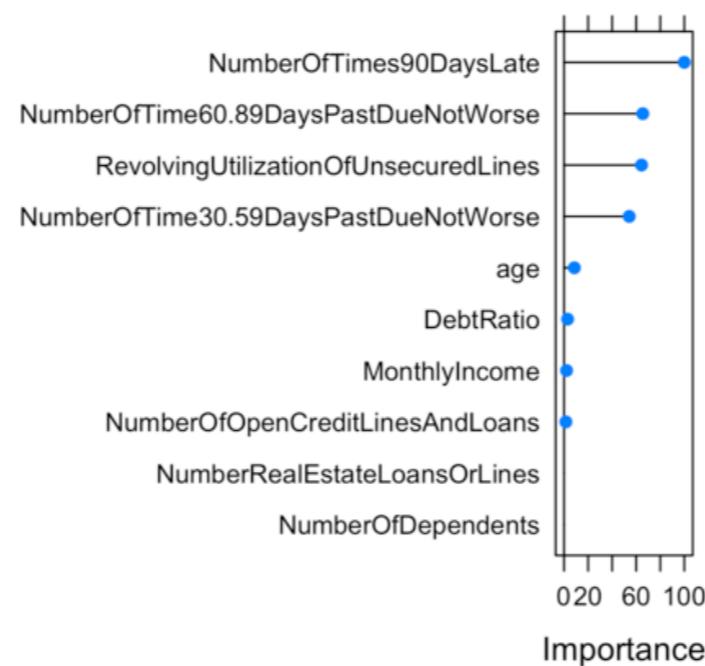


VARIABLE IMPORTANCE - OBS IN MODEL VS. PREDICTED

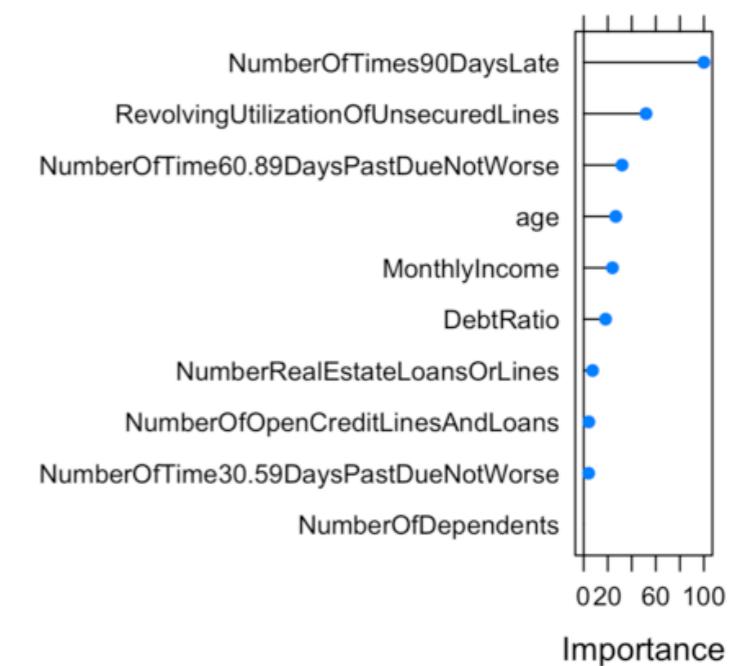
LogReg, 1st Iter.



CART, 1st Iter.

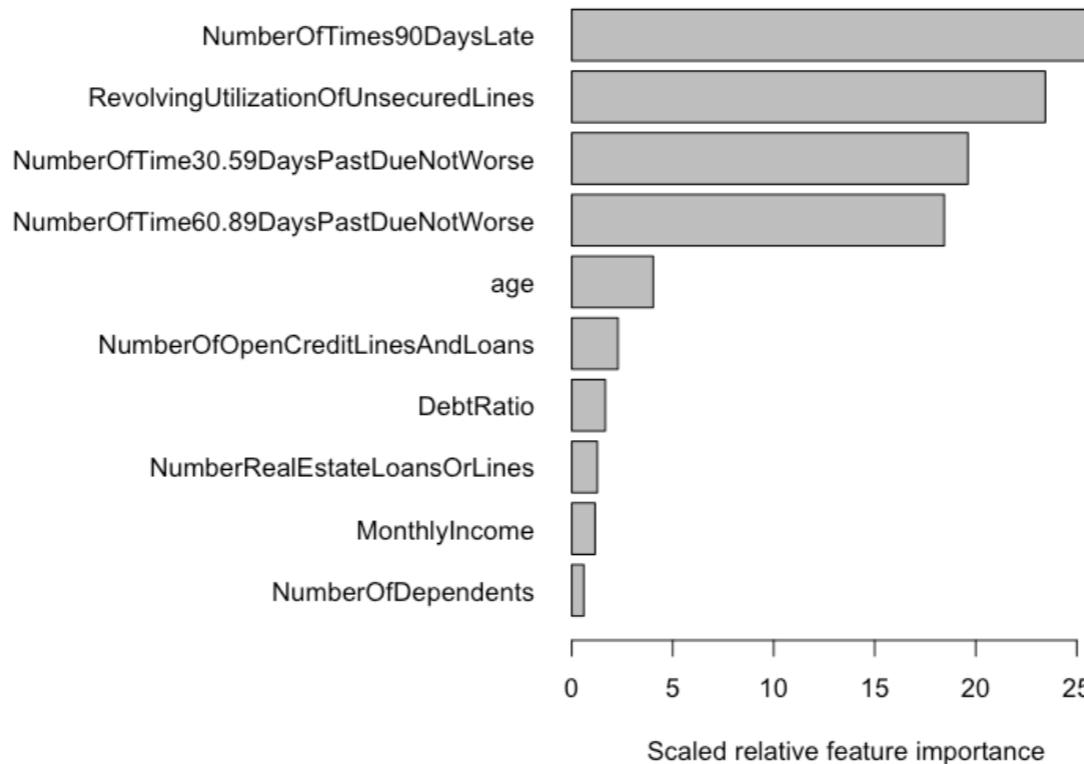


RF, 1st Iter.

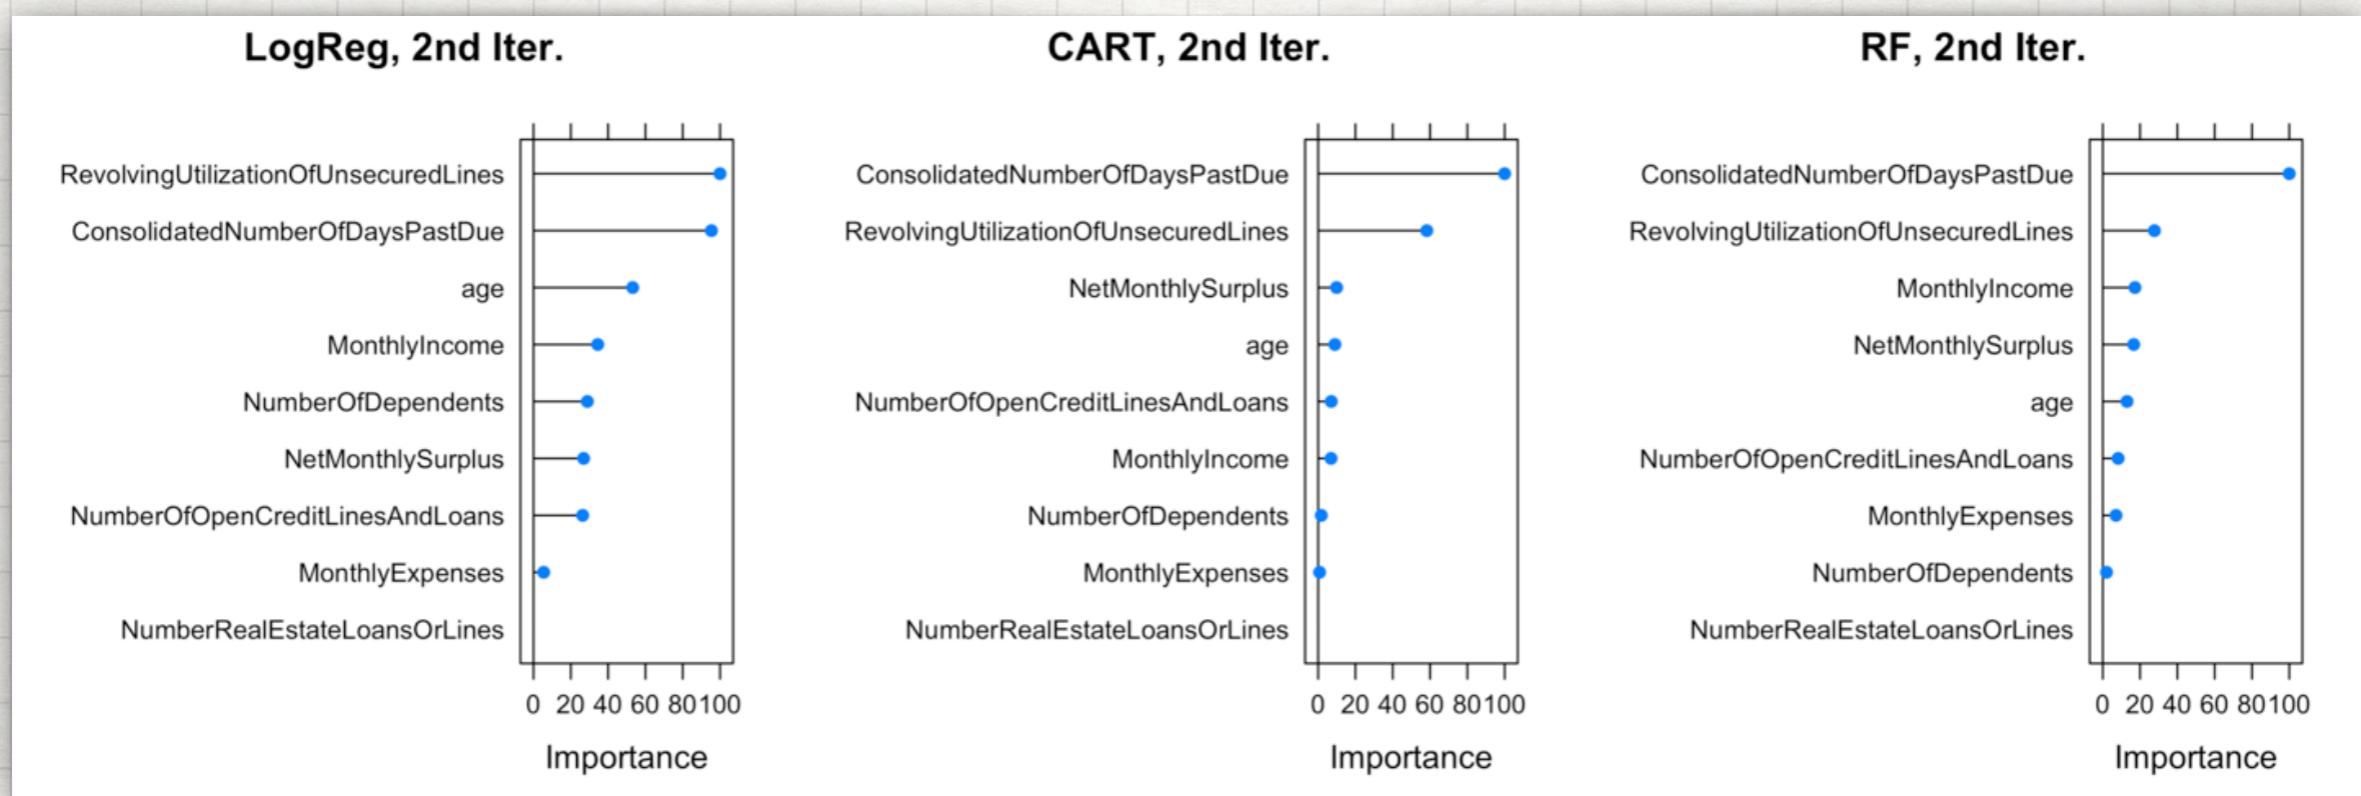


In First Iteration,
models built with
original variables.

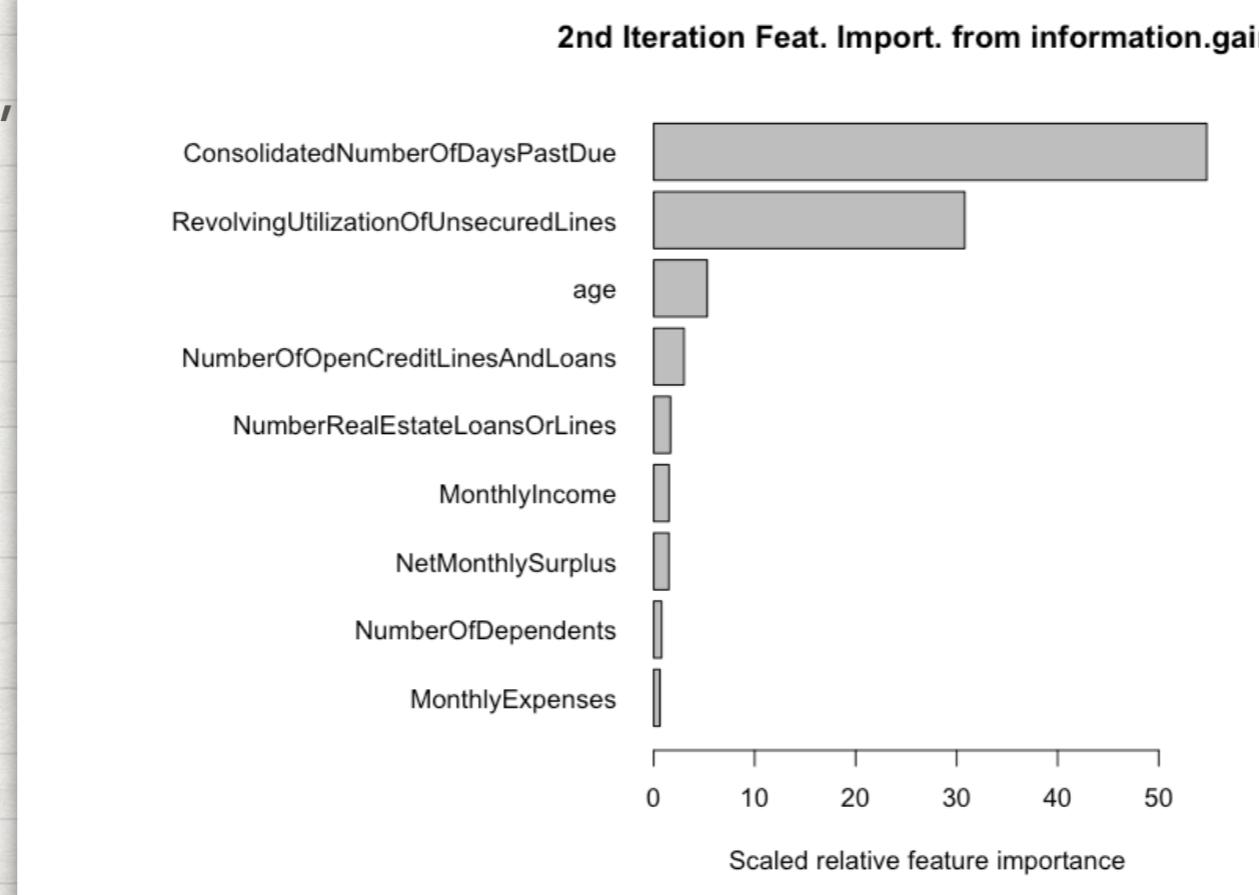
Original Feature Importance from information.gain()



VARIABLE IMPORTANCE - OBS IN MODEL VS. PREDICTED



In Second Iteration,
models built with
constructed
features.



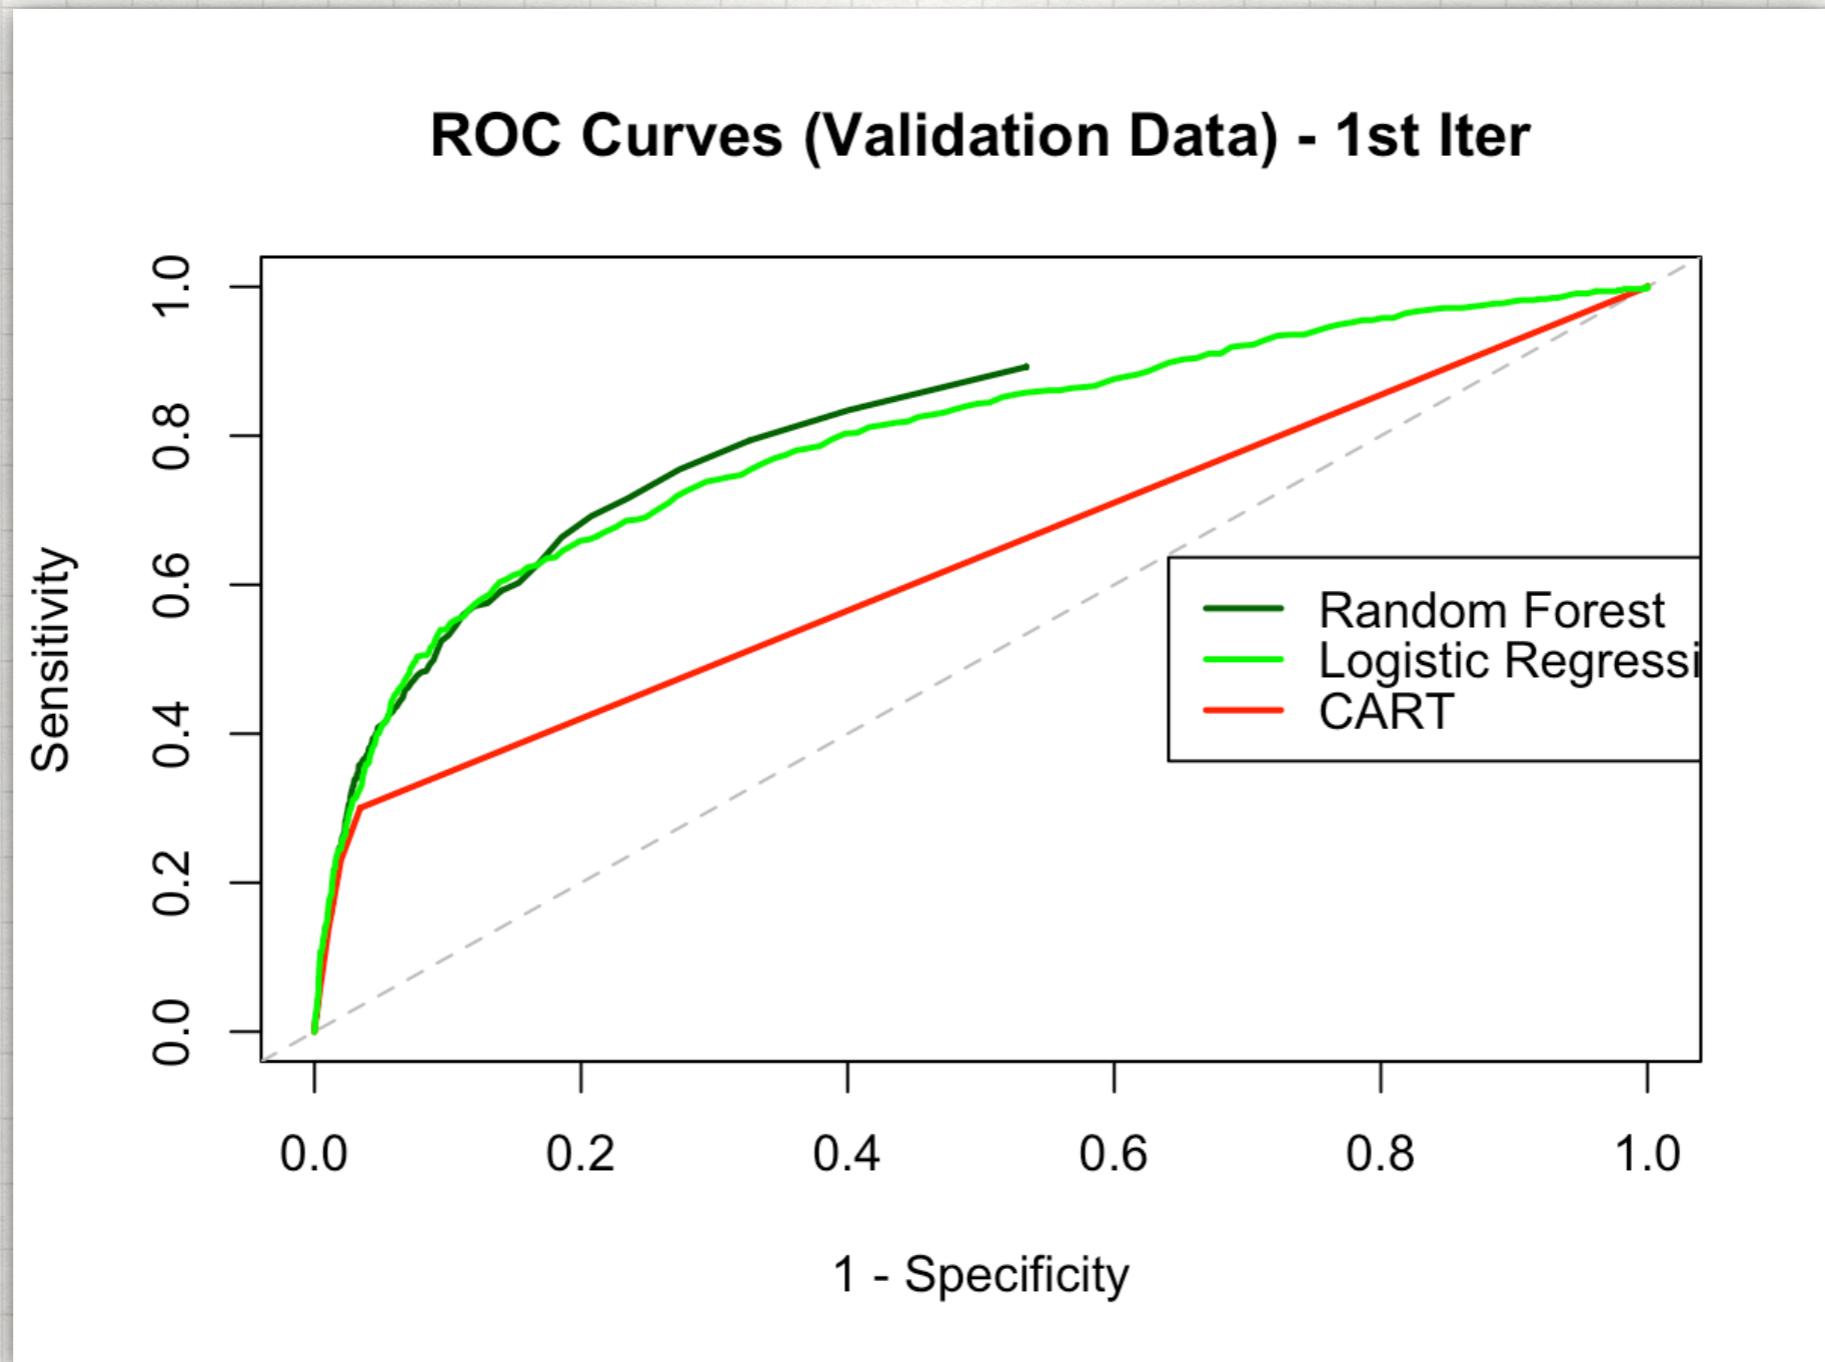
EVALUATE THE MODELS

COMPARE THE RELATIVE PREDICTIVE PERFORMANCE

- Mainly ROC Curves were used to evaluate models
- ROC curve visually demonstrates how much better model does at predicting the dependent binary variable compared to choosing an observation at random from the sample data.
- The area under the curve (AUC) on an ROC chart provides a quantitative measure of how well the model does over random predictions. Baseline performance is shown on ROC charts as a diagonal line from (0,0) to (1,1) and represents an AUC of 0.5-- which is the chance that a model will, at random, predict correctly the dependent variable.

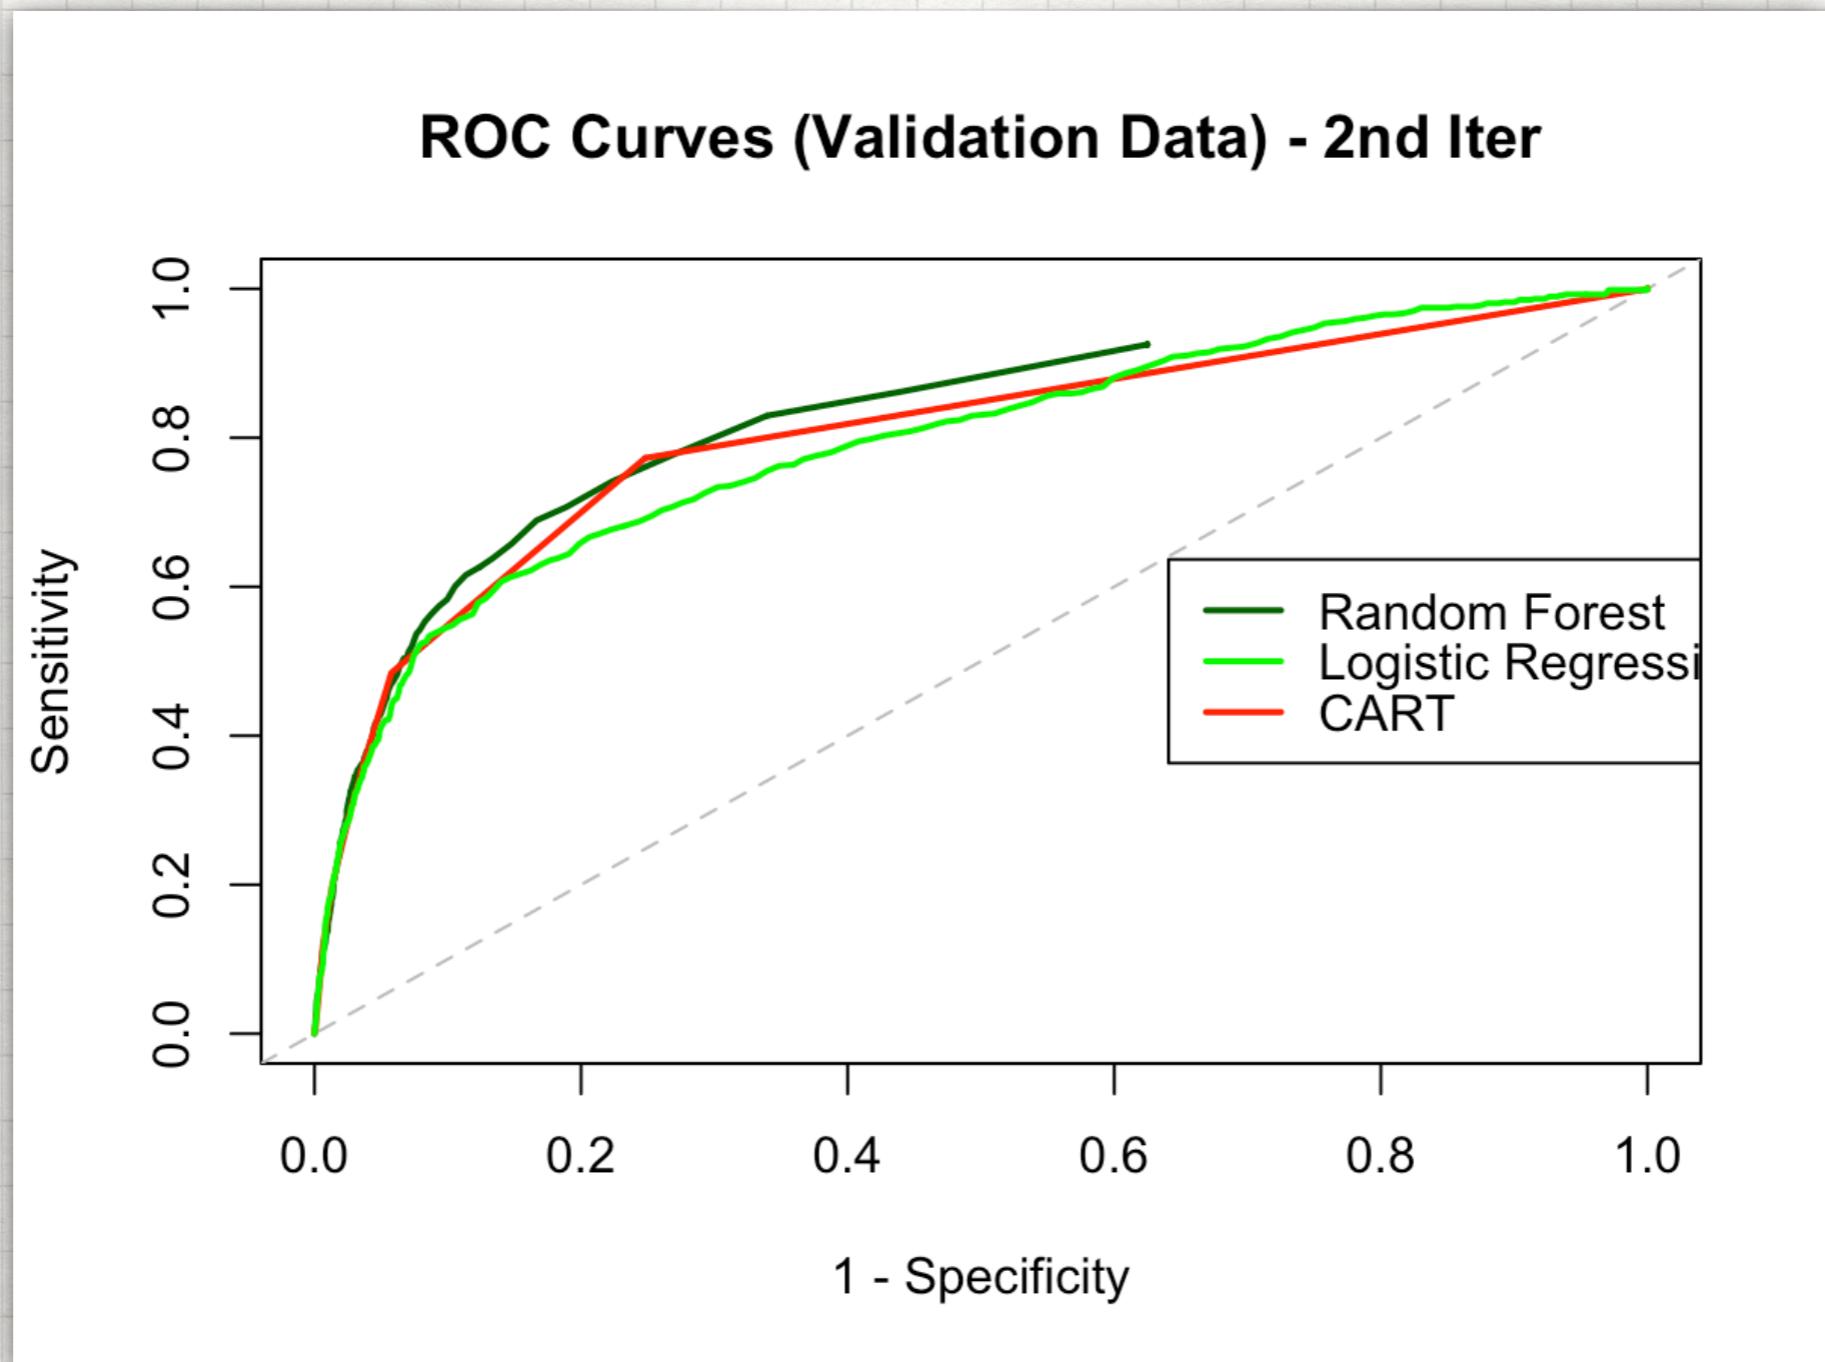
FIRST ITERATION

BUILT WITH ORIGINAL INDEPENDENT VARIABLES IN DATASET



SECOND ITERATION - FEATURE CHANGES

SOME CONSTRUCTED FEATURES USED IN MODEL BUILD



WHICH MODEL TYPES PERFORMED BEST?

OVER THE TWO ITERATIONS, THE ROC CURVES SHOW IT

- The **CART model** performed much better in second iteration with the engineered features included.
- The **logistic regression model** actually degraded slightly in second iteration
- The **Random Forest model** is clearly the best performing model among the three. It improved its lead in the second iteration.

KAGGLE SUBMISSION

CART, RF AND RF2

- Submitted a few entries on kaggle competition website
- As expected the RF (Random Forest) models performed the best. However, using the entire sample 150,000 training dataset was time consuming
- Below is how my best RF-based submission fared, better than only ~1/5 of teams that were in the competition, with an AUC of **0.832493**. But not bad for a novice like myself I think! My data science journey has started.

712	↑6	dataEngine	0.832583	1	Thu, 01 Dec 2011 03:44:52
-		DCrook	0.832493	-	Tue, 10 Nov 2015 00:51:28 Post-Deadline
Post-Deadline Entry If you would have submitted this entry during the competition, you would have been around here on the leaderboard.					
713	↑10	PZ	0.832043	8	Sun, 27 Nov 2011 02:44:46 (-24.4h)

KEY LEARNINGS

MAIN LESSONS I LEARNED

- Data cleanup is important and time consuming
- Feature engineering can help make better use of strangely defined data variables
- Obtaining better performing machine learning models can be computationally expensive. Ensemble model techniques can improve performance. Much to learn on this.
- Kaggle competitions really are competitive. They seem to be a good way to hone machine learning chops and learn about other state-of-the-art techniques.